

CarDekho Price Prediction

About Car Dehko

CarDekho is a prominent Indian automotive technology company headquartered in Jaipur, Rajasthan. Founded in 2008 by brothers Amit Jain and Anurag Jain, the platform assists users in buying and selling cars by providing comprehensive automotive content, including expert reviews, detailed specifications, pricing comparisons, and multimedia content for various car brands and models available in India.

Dataset Overview

The dataset contains the following columns:

car_name: Name of the car

brand: Car brand

model: Car model

vehicle_age: Age of the vehicle in years

km_driven: Total kilometers driven

seller_type: Type of seller (Individual or Dealer)

fuel_type: Type of fuel (Petrol, Diesel, etc.)

transmission_type: Manual or Automatic

mileage: Mileage of the car

engine: Engine capacity (in cc)

max_power: Maximum power output

seats: Number of seats

selling_price: Price at which the car was sold

Problem Statement:

The used car market in India is a dynamic and ever-changing landscape. Prices can fluctuate wildly based on a variety of factors including the make and model of the car, its mileage, its condition and the current market conditions. As a result, it can be difficult for sellers to accurately price their cars.

Approach:

We propose to develop a machine learning model that can predict the price of a used car based on its features. The model will be trained on a dataset of used cars that have been sold on

Cardekho.com in India. The model will then be able to be used to predict the price of any used car, given its features.

Objective

To build suitable Machine Learning Model for Used Car Price Prediction.

Benefits:

The benefits of this solution include:

Sellers will be able to more accurately price their cars which will help them to sell their cars faster and for a higher price.

Buyers will be able to find cars that are priced more competitively.

The overall used car market in India will become more efficient.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error,
r2_score

# Load dataset
df = pd.read_csv("Cardekho.csv")
df
```

	car_name	brand	model	vehicle_age	km_driven \
0	Maruti Alto	Maruti	Alto	9	120000
1	Hyundai Grand	Hyundai	Grand	5	20000
2	Hyundai i20	Hyundai	i20	11	60000
3	Maruti Alto	Maruti	Alto	9	37000
4	Ford Ecosport	Ford	Ecosport	6	30000
...
15406	Hyundai i10	Hyundai	i10	9	10723
15407	Maruti Ertiga	Maruti	Ertiga	2	18000
15408	Skoda Rapid	Skoda	Rapid	6	67000
15409	Mahindra XUV500	Mahindra	XUV500	5	3800000
15410	Honda City	Honda	City	2	13000

	seller_type	fuel_type	transmission_type	mileage	engine
max_power \					
0	Individual	Petrol	Manual	19.70	796
46.30					
1	Individual	Petrol	Manual	18.90	1197
82.00					
2	Individual	Petrol	Manual	17.00	1197

80.00						
3	Individual	Petrol	Manual	20.92	998	
67.10						
4	Dealer	Diesel	Manual	22.77	1498	
98.59						
...
...						
15406	Dealer	Petrol	Manual	19.81	1086	
68.05						
15407	Dealer	Petrol	Manual	17.50	1373	
91.10						
15408	Dealer	Diesel	Manual	21.14	1498	
103.52						
15409	Dealer	Diesel	Manual	16.00	2179	
140.00						
15410	Dealer	Petrol	Automatic	18.00	1497	
117.60						

	seats	selling_price
0	5	120000
1	5	550000
2	5	215000
3	5	226000
4	5	570000
...
15406	5	250000
15407	7	925000
15408	5	425000
15409	7	1225000
15410	5	1200000

[15411 rows x 13 columns]

```
# Display basic info
print(df.info())
print(df.head())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 15411 entries, 0 to 15410
Data columns (total 13 columns):
#   Column              Non-Null Count  Dtype
---  -
0   car_name            15411 non-null  object
1   brand               15411 non-null  object
2   model              15411 non-null  object
3   vehicle_age        15411 non-null  int64
4   km_driven           15411 non-null  int64
5   seller_type        15411 non-null  object
6   fuel_type           15411 non-null  object
7   transmission_type  15411 non-null  object
```

```

8  mileage      15411 non-null float64
9  engine       15411 non-null int64
10 max_power    15411 non-null float64
11 seats       15411 non-null int64
12 selling_price 15411 non-null int64

```

```
dtypes: float64(2), int64(5), object(6)
```

```
memory usage: 1.5+ MB
```

```
None
```

	car_name	brand	model	vehicle_age	km_driven
seller_type \					
0	Maruti Alto	Maruti	Alto	9	120000
Individual					
1	Hyundai Grand	Hyundai	Grand	5	20000
Individual					
2	Hyundai i20	Hyundai	i20	11	60000
Individual					
3	Maruti Alto	Maruti	Alto	9	37000
Individual					
4	Ford Ecosport	Ford	Ecosport	6	30000
Dealer					

	fuel_type	transmission_type	mileage	engine	max_power	seats	\
0	Petrol	Manual	19.70	796	46.30	5	
1	Petrol	Manual	18.90	1197	82.00	5	
2	Petrol	Manual	17.00	1197	80.00	5	
3	Petrol	Manual	20.92	998	67.10	5	
4	Diesel	Manual	22.77	1498	98.59	5	

```
selling_price
```

```
0      120000
```

```
1      550000
```

```
2      215000
```

```
3      226000
```

```
4      570000
```

```
# Handling missing values
```

```
df.dropna(inplace=True)
```

```
# Encoding categorical variables
```

```
label_encoders = {}
```

```
categorical_cols = df.select_dtypes(include=['object']).columns
```

```
for col in categorical_cols:
```

```
    le = LabelEncoder()
```

```
    df[col] = le.fit_transform(df[col])
```

```
    label_encoders[col] = le
```

```
# Splitting data
```

```
X = df.drop(columns=['selling_price']) # Assuming 'Price' is the
target column
```

```
y = df['selling_price']
```

```

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Scaling features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Model training
model = RandomForestRegressor(n_estimators=100, random_state=42)
model.fit(X_train, y_train)

RandomForestRegressor(random_state=42)

# Predictions
y_pred = model.predict(X_test)


# Evaluation
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred)

print(f'MAE: {mae}')
print(f'MSE: {mse}')
print(f'RMSE: {rmse}')
print(f'R2 Score: {r2}')

MAE: 99264.66682291472
MSE: 44451792219.313156
RMSE: 210835.9367359207
R2 Score: 0.9409499634804136

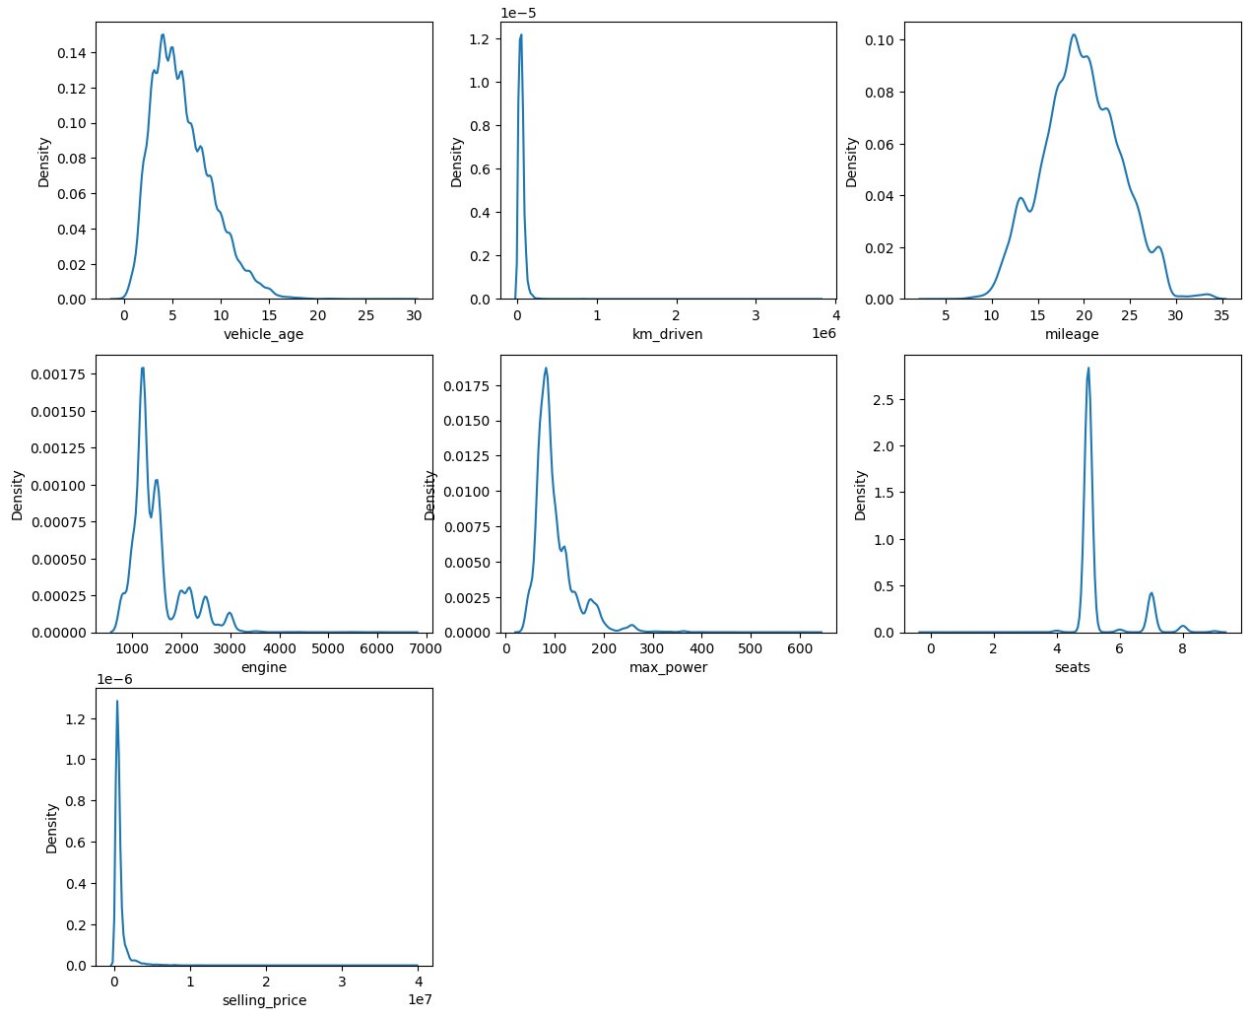
# Feature Importance
feature_importance = pd.Series(model.feature_importances_,
index=X.columns).sort_values(ascending=False)
plt.figure(figsize=(10, 6))
sns.barplot(x=feature_importance, y=feature_importance.index)
plt.xlabel('Feature Importance Score')
plt.ylabel('Features')
plt.title('Visualizing Important Features')
plt.show()

```



```
# Plot the density graphs of each of the numerical columns
num_columns = ['vehicle_age', 'km_driven', 'mileage', 'engine',
               'max_power', 'seats', 'selling_price']
plt.figure(figsize=(15, 30))
for i in range(len(num_columns)):
    plt.subplot(7, 3, i+1)
    sns.kdeplot(data = df[num_columns[i]])

plt.show()
```



Overall Insights

Right-Skewed Variables:

Variables like `km_driven`, `selling_price`, and `max_power` show strong right skewness, indicating a dominance of lower or typical values with a few outliers at the higher end. These variables may require transformation (e.g., log or square root) if used for modeling.

Concentrated Distributions:

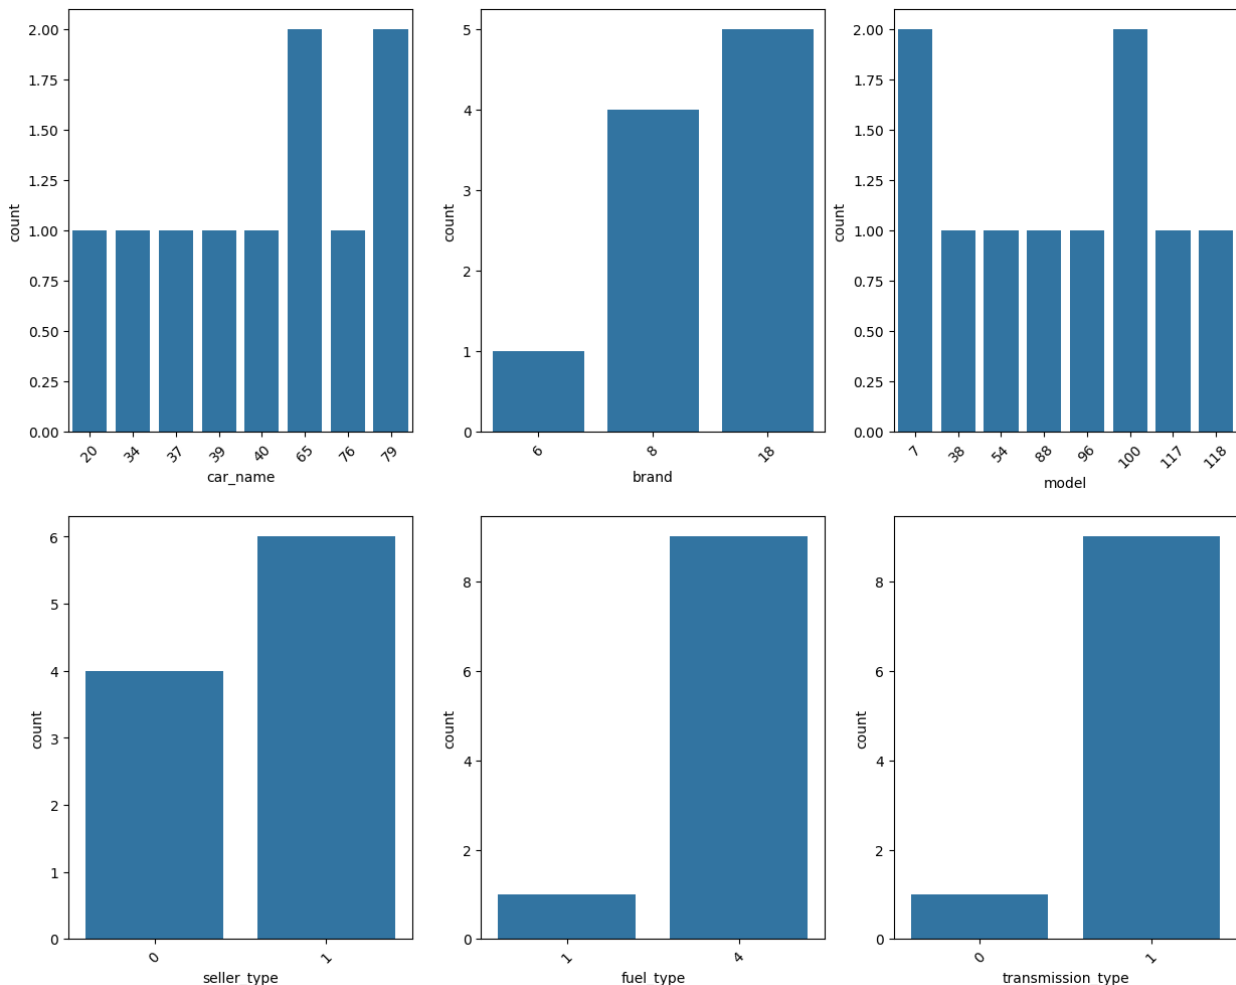
`Seats` and `vehicle_age` are highly concentrated around typical values (e.g., 5 seats, 0-10 years). This suggests standard consumer preferences for newer vehicles and compact/mid-sized cars.

Distinct Vehicle Segments:

The multimodal nature of `engine` and the broad range of `max_power` suggest that the dataset includes a mix of vehicle classes (e.g., compact, SUV, performance).

```
# Plot the countplots of each of the categorical columns
cat_columns = ['car_name', 'brand', 'model', 'seller_type',
               'fuel_type', 'transmission_type']
plt.figure(figsize=(15, 12))
for i in range(len(cat_columns)):
    plt.subplot(2, 3, i+1)
    plt.xticks(rotation = 45)
    sns.countplot(x = df[cat_columns[i]].head(10))

plt.show()
```



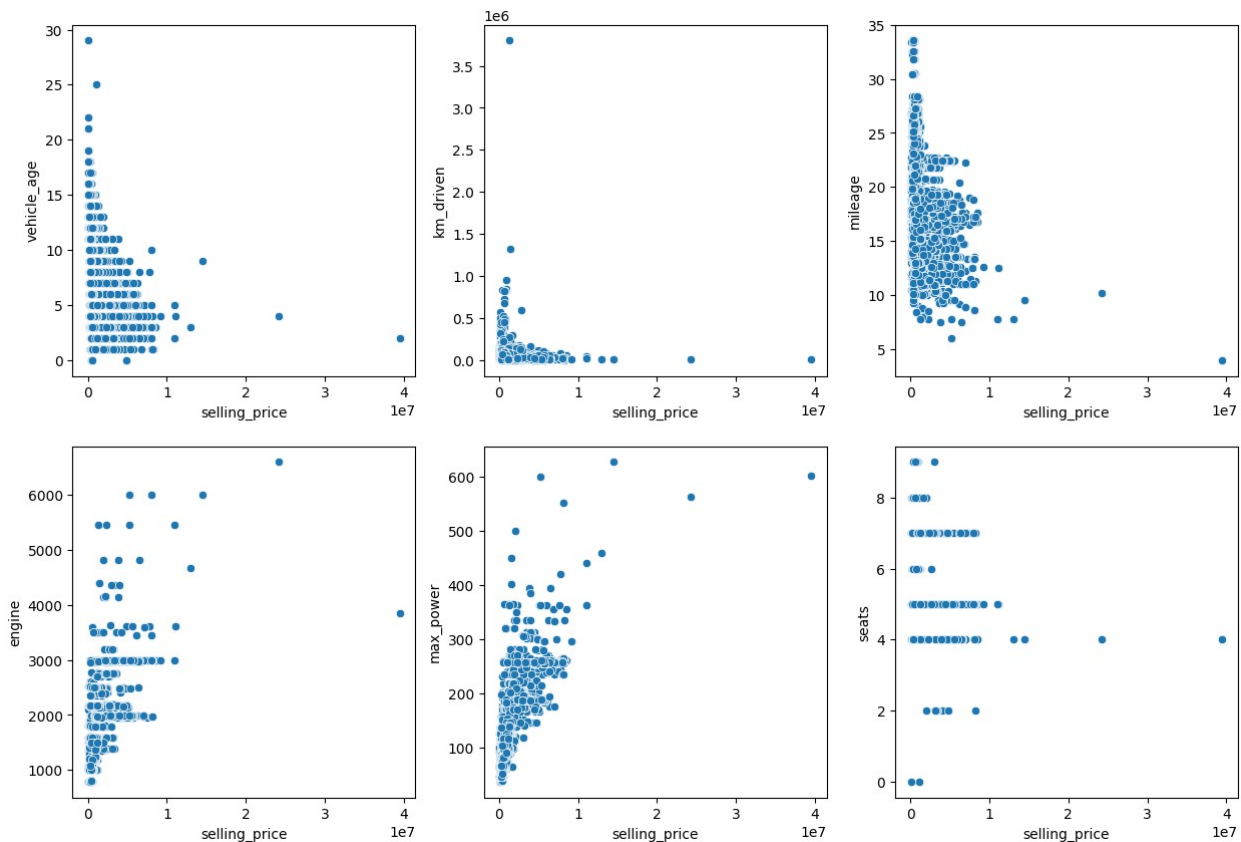
Insights

The dataset skews toward manual transmission and petrol-fueled cars sold by individuals. Maruti is the dominant brand, with the Maruti Alto and Hyundai Grand leading among models. Diesel cars and automatic transmissions are underrepresented, potentially indicating limited availability or demand in the dataset.


```
# Plot the relationship of each variable with the selling price
(Target variable)
numerical_columns = ['vehicle_age', 'km_driven', 'mileage', 'engine',
                    'max_power', 'seats']

plt.figure(figsize=(15, 10))
for i in range(len(numerical_columns)):
    plt.subplot(2, 3, i+1)
    sns.scatterplot(data = df, x = 'selling_price', y =
numerical_columns[i])

plt.show()
```



Interpretation of the Scatterplots

1. Vehicle Age vs Selling Price Observation: Older vehicles tend to have lower selling prices, indicating an inverse relationship.

Outliers: Some older vehicles show unusually high selling prices.

2. Kilometers Driven vs Selling Price Observation: Cars with fewer kilometers driven are priced higher, while cars with high mileage are clustered at lower prices.

Clusters: There's a noticeable cluster at low prices and low mileage.

3.*** Mileage vs Selling Price Observation:*** No strong correlation is visible. However, cars with lower mileage seem to cluster around average selling prices.

4. Engine vs Selling Price Observation: A positive relationship is visible: cars with larger engine capacities tend to have higher prices.

5. Max Power vs Selling Price Observation: A positive trend is visible: higher power correlates with higher selling prices.

Outliers: Some extreme outliers are present with very high power values.

6. Seats vs Selling Price Observation: No clear relationship is visible. Most cars have 4 or 5 seats, but prices don't vary significantly based on seating capacity.

